What Kubernetes doesn't do?

Kubernetes provides general features like deployment, scaling load balancing but also provides users with flexibility to choose whatever other pluggable solutions they would like to use for example for monitoring. It provides the mean to run services like middleware, databases, data-processing frameworks, etc. without providing these application-level services as built-in services. And, with use of portable mechanism these services can be accessed by other applications running on Kubernetes. Kubernetes focuses on orchestration and flexibility so, there aren't any mandatory solutions one need to use for logging, monitoring or alerting. It lets users integrate their logging, monitoring and alerting solutions by providing them with choices and flexibility along with the building blocks to build their platform. Kubernetes doesn't limit any types of applications supported, any application which can run in a container can run on Kubernetes. It doesn't deploy any source code, nor does it build your application. Due to use of a declarative API, there isn't any mandatory configuration language for Kubernetes. Kubernetes simply helps eliminate the need for orchestration by having independent, composable control processes which drive the state towards given desired state. The goal is to have an easier to use system which is more powerful, robust, resilient, and extensible.

References:

https://kubernetes.io/docs/concepts/overview/